

# Table of Contents

---

## Table of Contents

### Basic Definitions

- Disk
- Cryptography
  - Hashing
  - Encryption

### Simulation

- Definitions
  - LTS
  - LTS Refinement Relations
    - Refines to Related
    - Refines to Valid
    - Compiles to Valid
    - Transition Preserves Validity
  - Restrictions on Refinement Relations
    - High Oracle Exists
    - Oracle Refines to Same from Related
  - Self Simulation
  - Strong Bisimulation
- Metatheory
  - Main Theorem

### Layers

- Disk Layer (Layer 1)
  - Definitions
  - Operational Semantics
  - Key Cryptographic Assumptions

### Components

- Log
  - Structure
  - Header Contents
  - Functions
- Block Allocator
  - Implementation
    - Definitions
    - Representation Invariant
    - Functions
  - Abstraction (Block Allocator Layer)
    - Definitions
    - Operational Semantics
- Refinement

# Basic Definitions

---

## Disk

---

```
1  Definition addr := nat.
2  Axiom addr_eq_dec : forall (a b: addr), {a=b}+{a<>b}.
3
4  Axiom value : Type.
5  Axiom nat_to_value : nat -> value.
6  Axiom value_to_nat : value -> nat.
7
8  Axiom nat_to_value_to_nat:
9      forall n, value_to_nat (nat_to_value n) = n.
10 Axiom value_to_nat_to_value:
11     forall v, nat_to_value (value_to_nat v) = v.
12 Axiom value_dec: forall v v': value, {v=v'}+{v<>v'}.
```

# Cryptography

## Hashing

```
1 | Axiom hash : Type.
2 | Axiom hash0 : hash.
3 | Axiom hash_dec: forall h1 h2: hash, {h1 = h2}+{h1<>h2}.
4 | Axiom hash_function: hash -> value -> hash.
5 |
6 | Definition hashmap := mem hash hash_dec (hash * value).
7 |
8 | Fixpoint rolling_hash h vl :=
9 |   match vl with
10 |   | nil => h
11 |   | cons v vl' => rolling_hash (hash_function h v) vl'
12 |   end.
13 |
14 | Fixpoint rolling_hash_list h vl :=
15 |   match vl with
16 |   | nil => nil
17 |   | cons v vl' =>
18 |     let h' := hash_function h v in
19 |     cons h' (rolling_hash_list h' vl')
20 |   end.
21 |
22 | Fixpoint hash_and_pair h vl :=
23 |   match vl with
24 |   | nil => nil
25 |   | cons v vl' =>
26 |     let h' := hash_function h v in
27 |     cons (h, v) (hash_and_pair h' vl')
28 |   end.
```

## Encryption

```
1 | Axiom key: Type.
2 | Axiom key_dec: forall k1 k2: key, {k1 = k2}+{k1<>k2}.
3 | Axiom encrypt: key -> value -> value.
4 | Axiom encrypt_ext: forall k v v', encrypt k v = encrypt k v' -> v = v'.
5 |
6 | Definition encryptionmap := mem value value_dec (key * value).
```

# Simulation

## Definitions

### LTS

```
1 Record LTS :=
2   {
3     Oracle : Type;
4     State : Type;
5     Prog : Type -> Type;
6     exec : forall T, Oracle -> State -> Prog T -> @Result State T -> Prop;
7   }.
```

### LTS Refinement Relations

We need these relations because Coq doesn't have a good way to derive LTS's from existing ones by restricting state spaces, transition relations etc.

#### Refines to Related

This relation allow us to restrict properties to two low level states that refines to two high level states that are related by `related_h`.

Input:

- two input states (`s11` and `s12`),
- a refinement relation (`refines_to`), and
- a relation (`related_h`)

Assertions

- there are two other states (`sh1` and `sh2`) such that,
- `s11` (`s12`) refines to `sh1` (`sh2`) via `refines_to` relation, and
- `sh1` and `sh2` are related via `related_h`

```
1 Definition refines_to_related {State_L State_H: Type}
2   (refines_to: State_L -> State_H -> Prop)
3   (related_h: State_H -> State_H -> Prop)
4   (s11 s12: State_L)
5   : Prop :=
6   exists (sh1 sh2: State_H),
7     refines_to s11 sh1 /\
8     refines_to s12 sh2 /\
9     related_h sh1 sh2.
```

#### Refines to Valid

This definition allows us to restrict properties to low level states that refine to a valid high level state.

## Input

- an input state (`s1`),
- a refinement relation (`refines_to`), and
- a validity predicate (`valid_state_h`)

## Assertions

- for all states `sh`,
- if `s1` refines to `sh` via `refines_to` relation,
- then `sh` is a valid state (satisfies `valid_state_h`)

```
1 Definition refines_to_valid {State_L State_H: Type}
2   (refines_to: State_L -> State_H -> Prop)
3   (valid_state_h: State_H -> Prop)
4   (s1: State_L)
5   : Prop :=
6   forall (sh: State_H),
7     refines_to s1 sh ->
8     valid_state_h sh.
```

## Compiles to Valid

This definition allows us to restrict properties on low level programs that are a valid compilation of a high level program.

## Input

- an input program (`p1`),
- a refinement relation (`refines_to`), and
- a validity predicate (`valid_prog_h`)

## Assertions

- there is a program `ph` such that,
- `p1` is compilation of `ph`, and
- `ph` is a valid program (satisfies `valid_prog_h`)

```
1 Definition compiles_to_valid {Prog_L Prog_H: Type -> Type}
2   (valid_prog_h: forall T, Prog_H T -> Prop)
3   (compilation_of: forall T, Prog_L T -> Prog_H T -> Prop)
4   (T: Type)
5   (p1: Prog_L T)
6   : Prop :=
7   exists (ph: Prog_H T),
8     compilation_of T p1 ph /\
9     valid_prog_h T ph.
```

## Transition Preserves Validity

This definition ties notion of validity with restricting transitions by stating that any reachable state from a valid state is also valid.

```
1 Definition exec_compiled_preserves_validity lts1 lts2
2   (compilation_of: forall T, Prog lts1 T -> Prog lts2 T -> Prop)
  valid_state :=
3   forall T (p1: Prog lts1 T) (p2: Prog lts2 T) o s ret,
4     compilation_of T p1 p2 ->
5     valid_state s ->
6     exec lts1 T o s p1 ret ->
7     valid_state (extract_state ret).
```

## Restrictions on Refinement Relations

Following two properties ensures that your refinement relations has desired properties that allows transferring self simulations between layers

### High Oracle Exists

This definition states that

for all low level oracles `o1`

which results in a successful execution of a compiled program `p1` (that is compilation of `p2`)

from a low level state `s1` (that refines to a high level state),

there exists an high level oracle `o2` (that is a refinement of `o1`) for `p2`.

```
1 Definition high_oracle_exists {low high}
2   (refines_to: State low -> State high -> Prop)
3   (compilation_of : forall T, Prog low T -> Prog high T -> Prop)
4   (oracle_refines_to : forall T, State low -> Prog high T ->
Oracle low -> Oracle high -> Prop) :=
5   forall T o1 s1 s1' p1 p2,
6     (exists sh, refines_to s1 sh) ->
7     exec low T o1 s1 p1 s1' ->
8     compilation_of T p1 p2 ->
9     exists o2, oracle_refines_to T s1 p2 o1 o2.
```

### Oracle Refines to Same from Related

This definition states that our oracle refinement is agnostic to low level states that refine to related high level states. This property captures the fact that if two states are related, then they don't change the nondeterminism in different ways during refinement.

```
1 Definition oracle_refines_to_same_from_related {low high}
2   (refines_to: State low -> State high -> Prop)
3   (related_states_h: State high -> State high -> Prop)
4   (oracle_refines_to: forall T, State low -> Prog high T -> Oracle
low -> Oracle high -> Prop) :=
5   forall T o oh s1 s2 p2,
6     refines_to_related refines_to related_states_h s1 s2 ->
7     oracle_refines_to T s1 p2 o oh ->
8     oracle_refines_to T s2 p2 o oh.
```

## Self Simulation

This is a generalized two-safety property definition. Data confidentiality will be an instance of this. This is a little more stronger than a standard simulation because it forces two transitions in two executions to be the same.

```
1
2 Record SelfSimulation (lts: LTS)
3   (valid_state: State lts -> Prop)
4   (valid_prog: forall T, Prog lts T -> Prop)
5   (R: State lts -> State lts -> Prop) :=
6 {
7   self_simulation_correct:
8     forall T o p s1 s1' s2,
9       valid_state s1 ->
10      valid_state s2 ->
11      valid_prog T p ->
12      (exec lts) T o s1 p s1' ->
13      R s1 s2 ->
14      exists s2',
15        (exec lts) T o s2 p s2' /\
16        result_same s1' s2' /\
17        R (extract_state s1') (extract_state s2') /\
18        (forall def, extract_ret def s1' = extract_ret def s2') /\
19        valid_state (extract_state s1') /\
20        valid_state (extract_state s2') ;
21 }
```



## Strong Bisimulation

This is our refinement notion between two LTS's. It is stronger than a standard bisimulation because it requires transitions to be coupled, instead of just existing a transition in other LTS.

```
1 Record StrongBisimulation
2   (lts1 lts2 : LTS)
3   (compilation_of: forall T, Prog lts1 T -> Prog lts2 T -> Prop)
4   (refines_to: State lts1 -> State lts2 -> Prop)
5   (oracle_refines_to: forall T, State lts1 -> Prog lts2 T -> Oracle
lts1 -> Oracle lts2 -> Prop)
6   :=
7   {
8     strong_bisimulation_correct:
9       (forall T p1 (p2: Prog lts2 T) s1 s2 o1 o2,
10
11         refines_to s1 s2 ->
12         compilation_of T p1 p2 ->
13         oracle_refines_to T s1 p2 o1 o2 ->
14
15         (forall s1',
16           (exec lts1) T o1 s1 p1 s1' ->
17           exists s2',
18             (exec lts2) T o2 s2 p2 s2' /\
19             result_same s1' s2' /\
20             refines_to (extract_state s1') (extract_state s2') /\
21             (forall def, extract_ret def s1' = extract_ret def s2'))
22 /\
23
24         (forall s2',
25           (exec lts2) T o2 s2 p2 s2' ->
26           exists s1',
27             (exec lts1) T o1 s1 p1 s1' /\
28             result_same s1' s2' /\
29             refines_to (extract_state s1') (extract_state s2') /\
30             (forall def, extract_ret def s1' = extract_ret def s2'))))
31   }.
```

# Metatheory

---

## Main Theorem

```
1  Theorem transfer_high_to_low:
2    forall low high
3
4      related_states_h
5      refines_to
6      compilation_of
7      oracle_refines_to
8
9      valid_state_h
10     valid_prog_h,
11
12     SelfSimulation
13       high
14       valid_state_h
15       valid_prog_h
16       related_states_h ->
17
18     StrongBisimulation
19       low
20       high
21       compilation_of
22       refines_to
23       oracle_refines_to ->
24
25     high_oracle_exists refines_to compilation_of oracle_refines_to ->
26
27     oracle_refines_to_same_from_related refines_to related_states_h
28     oracle_refines_to ->
29
30     exec_compiled_preserves_validity
31     low
32     high
33     compilation_of
34     (refines_to_valid
35      refines_to
36      valid_state_h) ->
37
38     SelfSimulation
39       low
40       (refines_to_valid
41        refines_to
42        valid_state_h)
43       (compiles_to_valid
44        valid_prog_h
45        compilation_of)
46       (refines_to_related
47        refines_to
48        related_states_h).
```

# Layers

---

## Disk Layer (Layer 1)

---

### Definitions

```
1 Inductive token :=
2   | Key : key -> token
3   | Crash : token
4   | Cont : token.
5
6 Definition oracle := list token.
7
8 Definition state := (((list key * encryptionmap) * hashmap) * disk (value *
9   list value)).
10
11 Inductive prog : Type -> Type :=
12   | Read : addr -> prog value
13   | Write : addr -> value -> prog unit
14   | GetKey : list value -> prog key
15   | Hash : hash -> value -> prog hash
16   | Encrypt : key -> value -> prog value
17   | Decrypt : key -> value -> prog value
18   | Ret : forall T, T -> prog T
19   | Bind : forall T T', prog T -> (T -> prog T') -> prog T'.
```

# Operational Semantics

```
1 Definition consistent (m: mem A AEQ V) a v :=
2   m a = None \/ m a = Some v.
3
4 Fixpoint consistent_with_upds m al vl :=
5   match al, vl with
6   | nil, nil => True
7   | a::al', v::vl' =>
8     consistent m a v /\
9     consistent_with_upds (upd m a v) al' vl'
10  | _, _ => False
11  end.
```

```
1 Inductive exec : forall T, oracle -> state -> prog T -> @Result state T -
2   > Prop :=
3   ...
4
5   | ExecHash :
6     forall em hm d h v,
7       let hv := hash_function h v in
8       consistent hm hv (h, v) ->
9       exec [Cont] (em, hm, d) (Hash h v) (Finished (em, (upd hm hv (h,
10      v)), d) hv)
11
12   | ExecEncrypt :
13     forall kl em hm d k v,
14       let ev := encrypt k v in
15       consistent em ev (k, v) ->
16       exec [Cont] (kl, em, hm, d) (Encrypt k v) (Finished (kl, (upd em
17      ev (k, v)), hm, d) ev)
18
19   | ExecDecrypt :
20     forall kl em hm d ev k v,
21       ev = encrypt k v ->
22       em ev = Some (k, v) ->
23       exec [Cont] (kl, em, hm, d) (Decrypt k ev) (Finished (kl, em, hm,
24      d) v)
25
26   | ExecGetKey :
27     forall vl kl em hm d k,
28       ~In k kl ->
29       consistent_with_upds em
30       (map (encrypt k) vl) (map (fun v => (k, v)) vl) ->
31       exec [Key k] (kl, em, hm, d) (GetKey vl) (Finished ((k::kl), em,
32      hm, d) k).
```

## Key Cryptographic Assumptions

### No Hash Collisions

This assumption embodied as execution getting stuck if a collision happens during execution. Each hashed value is stored in a map after execution and each input is checked before executing the hash operation.

#### *Justification*

It is exponentially unlikely to have a hash collision in a real system.

### No Encryption Collisions

This assumption embodied as execution getting stuck if a collision happens during execution. Each key and block pair is stored in a map and each input checked before executing the encryption operation. This is a stronger assumption than the traditional one because it requires no collision for (key, value) pairs instead of two values for the same key.

#### *Why do we need a stronger assumption?*

It is required to prevent the following scenario:

There are two equivalent states `st1` and `st2`.

- a transaction is committed,
- header and all but one data block makes to the disk
- crash happens
- non-written data block matches what is on the disk on `st1` only
- recovery commits in `st1` but not in `st2`, leaking confidential information.

#### *Justification*

In real execution, it is exponentially unlikely to have a collision even for (key, value) pairs. Also, even in the case that such collision happens, leaked data is practically garbage because it is encrypted.

### Generated Key Does Not Cause Collision

This strong assumption (combined with the total correctness requirement) enforces some restrictions for key generation in operational semantics. We need to ensure that generated key will not create an encryption collision. To ensure that, `GenKey` operation takes the blocks that will be encrypted as input.

#### *Justification*

In real execution, it is exponentially unlikely to have a collision. Also, even in the case that such collision happens, leaked data is practically garbage because it is encrypted.

One way to circumvent this would be combining `GenKey` and `Encrypt` operation into one operation

```
EncryptWithNewKey: list block -> prog (key * list block)
```

Which takes blocks to be encrypted and encrypts them with a new key, returning both key and the encrypted blocks. This operations limitation is not a problem for us because every time we are encrypting, we do it with a fresh key anyway.

# Components

---

## Log

---

### Structure

1	-----				
2			TXN 1	TXN 2	
3		Header	Addr	Data	. . .
4			Blocks	Blocks	
5	-----				

### Header Contents

```
1  Record txn_record :=
2    {
3      txn_key : key;
4      txn_start : nat; (* Relative to start of the log *)
5      addr_count : nat;
6      data_count : nat;
7    }.
8
9  Record header :=
10   {
11     old_hash : hash;
12     old_count : nat;
13     old_txn_count : nat;
14     cur_hash : hash;
15     cur_count : nat;
16     txn_records : list txn_record;
17   }.
```

## Functions

```
1 Definition commit (addr_l data_l: list value) :=
2   hdr <- read_header;
3   if (new_count <=? log_length) then
4     new_key <- GetKey (addr_l++data_l);
5     enc_data <- encrypt_all new_key (addr_l ++ data_l);
6     _ <- write_consecutive (log_start + cur_count) enc_data;
7     new_hash <- hash_all cur_hash enc_data;
8     _ <- write_header new_hdr;
9     Ret true
10  else
11    Ret false.
```

```
1 Definition apply_log :=
2   hdr <- read_header;
3   log <- read_consecutive log_start cur_count;
4   success <- check_and_flush txns log cur_hash;
5   if success then
6     Ret true
7   else
8     success <- check_and_flush old_txns old_log old_hash;
9     Ret success.
```

```
1 Definition check_and_flush txns log hash :=
2   log_hash <- hash_all hash0 log;
3   if (hash_dec log_hash hash) then
4     _ <- flush_txns txns log;
5     Ret true
6   else
7     Ret false.
8
9 Definition flush_txns txns log_blocks :=
10  _ <- apply_txns txns log_blocks;
11  _ <- write_header header0;
12  Ret tt.
13
14 Fixpoint apply_txns txns log_blocks :=
15  match txns with
16  | nil =>
17    Ret tt
18  | txn::txns' =>
19    _ <- apply_txn txn log_blocks;
20    _ <- apply_txns txns' log_blocks;
21    Ret tt
22  end.
23
24 Definition apply_txn txn log_blocks :=
25  plain_blocks <- decrypt_all key txn_blocks;
26  _ <- write_batch addr_list data_blocks;
27  Ret tt.
```



# Block Allocator

## Implementation

### Definitions

```
1 | Axiom block_size: nat.
2 | Axiom block_size_eq: block_size = 64.
3 |
4 | Definition valid_bitlist l :=
5 |   length l = block_size /\ (forall i, In i l -> i < 2).
6 |
7 | Record bitlist :=
8 | {
9 |   bits : list nat;
10 |   valid : valid_bitlist bits
11 | }.
12 |
13 | Axiom value_to_bits: value -> bitlist.
14 | Axiom bits_to_value: bitlist -> value.
15 | Axiom value_to_bits_to_value :
16 |   forall v, bits_to_value (value_to_bits v) = v.
17 | Axiom bits_to_value_to_bits :
18 |   forall l, value_to_bits (bits_to_value l) = l.
```

### Representation Invariant

```
1 | Definition rep (dh: disk value) : @pred addr addr_dec (set value) :=
2 |   (exists bitmap bl,
3 |     let bits := bits (value_to_bits bitmap) in
4 |     0 |-> (bitmap, bl) *
5 |     ptsto_bits dh bits *
6 |     [[ forall i, i >= block_size -> dh i = None ]])%pred.
```

`ptsto_bits dh bits` says that `forall i < length bits, bits[i] = 0 <-> dh (S i) = None`.

### Functions

```
1 | Definition read a : prog (option value) :=
2 |   v <- Read 0;
3 |   if a < block_size then
4 |     if bitmap[a] = 1 then
5 |       h <- Read (S a);
6 |       Ret (Some h)
7 |     else
8 |       Ret None
9 |   else
10 |     Ret None.
```

```

1 Definition write a v' : prog (option unit) :=
2   v <- Read 0;
3   if a < block_size then
4     if bitmap[a] = 1 then
5       Write (S a) v'
6     else
7       Ret None
8   else
9     Ret None.

```

```

1 Definition alloc (v': value) : prog (option addr) :=
2   v <- Read 0;
3   if first_zero bitmap < block_size then
4     _ <- Write (S (first_zero bitmap)) v';
5     _ <- Write 0 (to_block (updN bitmap (first_zero bitmap) 1));
6     Ret (Some index)
7   else
8     Ret None.

```

```

1 Definition free a : prog unit :=
2   v <- Read 0;
3   if a < block_size then
4     if bitmap[a] = 1 then
5       Write 0 (to_block (updN bitmap a 0));
6     else
7       Ret tt
8   else
9     Ret tt.

```

## Abstraction (Block Allocator Layer)

### Definitions

```

1 Inductive token:=
2   | BlockNum : addr -> token
3   | DiskFull : token
4   | Cont : token
5   | Crash1 : token
6   | Crash2 : token
7   | CrashAlloc: addr -> token.
8
9 Definition state := disk value.
10
11 Inductive prog : Type -> Type :=
12   | Read : addr -> prog (option value)
13   | Write : addr -> value -> prog (option unit)
14   | Alloc : value -> prog (option addr)
15   | Free : addr -> prog unit
16   | Ret : forall T, T -> prog T
17   | Bind : forall T T', prog T -> (T -> prog T') -> prog T'.

```

We need two different crash tokens because implementation and abstraction must have the same overall behavior (`Finished -> Finished` and `Crashed -> Crashed`). This prevents us from letting abstraction succeed if implementation crashes after "its commit point" even though resulting states may be the same. Therefore, in case of a crash, we approximately need to know where crash happened so that abstraction can follow accordingly.

## Operational Semantics

```

1 Inductive exec :
2   forall T, oracle -> state -> prog T -> @Result state T -> Prop :=
3
4   ...
5
6   | ExecAllocSucc :
7     forall d a v,
8       d a = None ->
9       exec [BlockNum a] d (Alloc v) (Finished (upd d a v) (Some a))
10
11  | ExecAllocFail :
12    forall d v,
13      exec [DiskFull] d (Alloc v) (Finished d None)
14
15  ...
16
17  | ExecAllocCrashBefore :
18    forall d v,
19      exec [Crash1] d (Alloc v) (Crashed d)
20
21  | ExecAllocCrashAfter :
22    forall d a v,
23      d a = None ->
24      exec [CrashAlloc a] d (Alloc v) (Crashed (upd d a v))
25
26  ...

```

## Refinement

```

1 Fixpoint oracle_refines_to T
2   (d1: State layer1_lts) (p: Layer2.prog T)
3   (o1: Oracle layer1_lts) (o2: Layer2.oracle) : Prop :=
4   oracle_ok (compile p) o1 d1 /\
5   match p with
6   | Alloc v =>
7     if In Crash o1 then
8       forall d1',
9         Layer1.exec o1 d1 (compile p) (Crashed d1') ->
10        (d1 0 = d1' 0 -> o2 = [Crash1]) /\
11        (d1' 0 <> d1' 0 -> o2 = [CrashAlloc first_zero])
12     else
13       if first_zero bitmap < block_size then
14         o2 = [BlockNum (first_zero bitmap)]
15       else

```

```

16         o2 = [DiskFull]
17     end
18 | Bind p1 p2 =>
19     exists o1' o1'',
20     o1 = o1' ++ o1'' /\
21     ((exists d1',
22         Layer1.exec o1 d1 (compile p1) (Crashed d1') /\
23         oracle_refines_to d1 p1 o1 o2 /\ o1'' = [])
24     /\
25     (exists d1' r ret,
26         Layer1.exec o1' d1 (compile p1) (Finished d1' r) /\
27         Layer1.exec o1'' d1' (compile (p2 r)) ret /\
28         (exists o2' o2'',
29         oracle_refines_to d1 p1 o1' o2' /\
30         oracle_refines_to d1' (p2 r) o1'' o2'' /\
31         o2 = o2' ++ o2''))
32     ...
33 end.
34
35 Definition refines_to d1 d2 :=
36     exists F, (F * rep d2)%pred d1.

```